

Deep Learning Pipelines on Databricks



databricks

CS570 Big Data Processing Project
By Feven Araya
Instructor: Dr. Chang, Henry

Table of contents

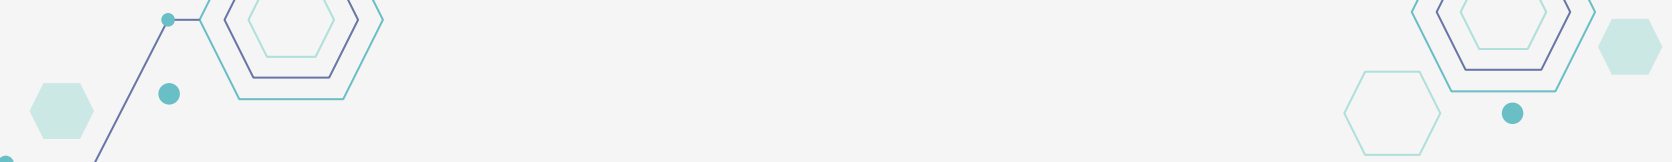
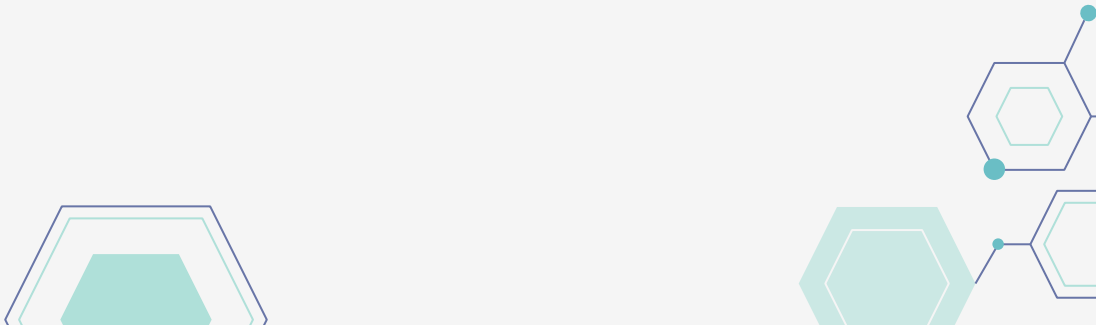
1. Introduction
2. Design
3. Implementation
4. Testing
5. Enhancement
6. Conclusion
7. References



01

Introduction




- 
- Deep Learning Pipelines is a new library from Databricks that offers high-level APIs to seamlessly integrate deep learning model application and transfer learning within Apache Spark's MLlib Pipelines and Spark SQL.
 - This library leverages popular deep learning libraries to enable scalable deep learning model application.
 - It allows users to work with images natively in Spark DataFrames, perform transfer learning, apply deep learning models at scale, and deploy models as SQL functions.
- 


02

Design






The design of Deep Learning Pipelines focuses on integrating deep learning capabilities within the Spark ecosystem. It provides tools for:

1. **Working with images in Spark DataFrames:** Utility functions load millions of images into a Spark DataFrame and decode them automatically in a distributed fashion.
 2. **Transfer Learning:** Facilitates leveraging existing deep learning models with minimal code and run-time effort.
 3. **Applying Deep Learning Models at Scale:** Uses Spark MLlib Transformers to apply TensorFlow Graphs and Keras Models efficiently.
 4. **Future Enhancements:** Plans to include distributed hyper-parameter tuning via Spark MLlib Pipelines and deploying models as SQL functions.
- 

03

Implementation



Cluster Setup

- 1. **Library Installation:** Deep Learning Pipelines is available as a Spark Package. Users need to create a new library with the Maven Coordinate source option to find "spark-deep-learning" and attach it to their cluster.
- 2. **Dependencies:** The notebook requires the following libraries via PyPI: tensorflow, keras, h5py.

```
1/22/2018 (4s) 6

%sh
curl -O http://download.tensorflow.org/example_images/flower_photos.tgz
tar xzf flower_photos.tgz
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current	
			Dload	Upload	Total	Spent	Left	Speed
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
12	218M	12	28.0M	0	0	24.0M	0	0:00:09 0:00:01 0:00:08 24.0M
100	218M	100	218M	0	0	114M	0	0:00:01 0:00:01 0:00:01 114M

```
tar: flower_photos/roses/14810868100_87eb739f26_m.jpg: Cannot change ownership to uid 270850, gid 5000: Invalid argument
tar: flower_photos/roses/1446090416_f0cad5fde4.jpg: Cannot change ownership to uid 270850, gid 5000: Invalid argument
tar: flower_photos/roses/15319767030_e6c5602a77_m.jpg: Cannot change ownership to uid 270850, gid 5000: Invalid argument
tar: flower_photos/roses/15032112248_30c5284e54_n.jpg: Cannot change ownership to uid 270850, gid 5000: Invalid argument
```

1/22/2018 (4s) 7

display(dbutils.fs.ls('file:/databricks/driver/flower_photos'))

Table

	path	name	size
1	file:/databricks/driver/flower_photos/daisy/	daisy/	32768
2	file:/databricks/driver/flower_photos/dandelion/	dandelion/	49152
3	file:/databricks/driver/flower_photos/sunflowers/	sunflowers/	36864
4	file:/databricks/driver/flower_photos/LICENSE.txt	LICENSE.txt	418049
5	file:/databricks/driver/flower_photos/tulips/	tulips/	40960
6	file:/databricks/driver/flower_photos/roses/	roses/	36864

Loading and Processing Images

- 1. **Loading Images:** The `readImages` function loads images into a Spark DataFrame.
- 2. **Creating Sample Set:** A smaller sample set is created for quick demonstrations by copying a subset of images.

```
# The 'file:/...' directory will be cleared out upon cluster termination. That doesn't matter for
images in a more permanent place. Let's move the files to dbfs so we can see how to work with it
img_dir = '/tmp/flower_photos'
dbutils.fs.mkdirs(img_dir)
dbutils.fs.cp('file:/databricks/driver/flower_photos/tulips', img_dir + "/tulips", recurse=True)
dbutils.fs.cp('file:/databricks/driver/flower_photos/daisy', img_dir + "/daisy", recurse=True)
dbutils.fs.cp('file:/databricks/driver/flower_photos/LICENSE.txt', img_dir)
display(dbutils.fs.ls(img_dir))
```

Table			
	path	name	size
1	dbfs:/tmp/flower_photos/LICENSE.txt	LICENSE.txt	418049
2	dbfs:/tmp/flower_photos/daisy/	daisy/	0
3	dbfs:/tmp/flower_photos/tulips/	tulips/	0

```
# Let's create a small sample set of images for quick demonstrations.
sample_img_dir = img_dir + "/sample"
dbutils.fs.mkdirs(sample_img_dir)
files = dbutils.fs.ls(img_dir + "/tulips") [0:1] + dbutils.fs.ls(img_dir + "/daisy") [0:2]
for f in files:
    dbutils.fs.cp(f.path, sample_img_dir)
display(dbutils.fs.ls(sample_img_dir))
```

Table			
	path	name	size
1	dbfs:/tmp/flower_photos/sample/100080576_f52e8ee070_n.jpg	100080576_f52e8ee070_n.jpg	26797
2	dbfs:/tmp/flower_photos/sample/100930342_92e8746431_n.jpg	100930342_92e8746431_n.jpg	26200
3	dbfs:/tmp/flower_photos/sample/10140303196_b88d3d6cec.jpg	10140303196_b88d3d6cec.jpg	117247

```
from sparkdl import readImages
image_df = readImages(sample_img_dir)
```

▶  image_df: pyspark.sql.dataframe.DataFrame = [filePath: string, image: struct]

```
/databricks/python/local/lib/python2.7/site-packages/h5py/__init__.py:36: Fut
is deprecated. In future, it will be treated as `np.float64 == np.dtype(float
from ._conv import register_converters as _register_converters
Using TensorFlow backend.
```

```
display(image_df)
```

Table ▼ +

	 filePath	 image
1	dbfs:/tmp/flower_photos/sample/100080576_f52e8ee070_n.jpg	> {"mode": "RGB", "height": 263, "width": 320, "nChannels": 3, "data": "h4eFioqljo6OkZGRkpKSk5OTIZWXl5eZmZmZl5eVIZWTIZW..."}
2	dbfs:/tmp/flower_photos/sample/100930342_92e8746431_n.jpg	> {"mode": "RGB", "height": 209, "width": 320, "nChannels": 3, "data": "Ey4PEC8QDi8QDi8SES4SEy0SDy0RDS4RCiYPCCQNByMN..."}

Transfer Learning

1. **Data Preparation:** Images are labeled and split into training and test DataFrames.
2. **Model Training:** A logistic regression model is trained using features extracted by the `DeepImageFeaturizer` from the InceptionV3 model.
3. **Model Evaluation:** The trained model's accuracy is evaluated on the test set.

```
1/22/2018 (1s) 15 Python

# Create training & test DataFrames for transfer learning - this piece of code is longer than transfer learning itself below!
from sparkdl import readImages
from pyspark.sql.functions import lit

tulips_df = readImages(img_dir + "/tulips").withColumn("label", lit(1))
daisy_df = readImages(img_dir + "/daisy").withColumn("label", lit(0))
tulips_train, tulips_test, _ = tulips_df.randomSplit([0.05, 0.05, 0.9]) # use larger training sets (e.g. [0.6, 0.4] for non-community edition clusters)
daisy_train, daisy_test, _ = daisy_df.randomSplit([0.05, 0.05, 0.9]) # use larger training sets (e.g. [0.6, 0.4] for non-community edition clusters)
train_df = tulips_train.unionAll(daisy_train)
test_df = tulips_test.unionAll(daisy_test)

# Under the hood, each of the partitions is fully loaded in memory, which may be expensive.
# This ensure that each of the partitions has a small size.
train_df = train_df.repartition(100)
test_df = test_df.repartition(100)
```



```
from pyspark.ml.classification import LogisticRegression
from pyspark.ml import Pipeline
from sparkdl import DeepImageFeaturizer
```

```
featurizer = DeepImageFeaturizer(inputCol="image", outputCol="features", modelName="InceptionV3")
lr = LogisticRegression(maxIter=20, regParam=0.05, elasticNetParam=0.3, labelCol="label")
p = Pipeline(stages=[featurizer, lr])
```

```
p_model = p.fit(train_df)
```

```
from pyspark.ml.evaluation import MulticlassClassificationEvaluator
```

```
tested_df = p_model.transform(test_df)
evaluator = MulticlassClassificationEvaluator(metricName="accuracy")
print("Test set accuracy = " + str(evaluator.evaluate(tested_df.select("prediction", "label"))))
```

▶ tested_df: pyspark.sql.dataframe.DataFrame

```
INFO:tensorflow:Froze 376 variables.
Converted 376 variables to const ops.
INFO:tensorflow:Froze 0 variables.
Converted 0 variables to const ops.
Test set accuracy = 0.971014492754
```



```

from pyspark.sql.types import DoubleType
from pyspark.sql.functions import expr
def _p1(v):
    return float(v.array[1])
p1 = udf(_p1, DoubleType())

df = tested_df.withColumn("p_1", p1(tested_df.probability))
wrong_df = df.orderBy(expr("abs(p_1 - label)", ascending=False))
display(wrong_df.select("filePath", "p_1", "label").limit(10))

```

- df: pyspark.sql.dataframe.DataFrame
- wrong_df: pyspark.sql.dataframe.DataFrame

Table ▾ +

	A_C^B filePath	1.2 p_1	1_3^2 label	
1	dbfs:/tmp/flower_photos/daisy/9345273630_af3550031d.jpg	0.804202936186138	0	
2	dbfs:/tmp/flower_photos/daisy/530738000_4df7e4786b.jpg	0.6165413243876156	0	
3	dbfs:/tmp/flower_photos/tulips/113902743_8f537f769b_n.jpg	0.5153939149058596	1	

Applying Deep Learning Models at Scale

1. **Using Pre-trained Models:** The `DeepImagePredictor` applies pre-trained models like InceptionV3 to images and returns predictions.
2. **Custom TensorFlow Graphs:** The `TFImageTransformer` applies a user-defined TensorFlow Graph to the image DataFrame.
3. **Keras Models:** The `KerasImageFileTransformer` loads a Keras model and applies it to the DataFrame containing image URIs.

```
from sparkdl import readImages, DeepImagePredictor

image_df = readImages(sample_img_dir)

predictor = DeepImagePredictor(inputCol="image", outputCol="predicted_labels", modelName="InceptionV3", decodePredictions=True, topK=10)
predictions_df = predictor.transform(image_df)

display(predictions_df.select("filePath", "predicted_labels"))
```

image_df: pyspark.sql.dataframe.DataFrame = [filePath: string, image: struct]
predictions_df: pyspark.sql.dataframe.DataFrame = [filePath: string, image: struct ... 1 more field]

Table ▾ +

	^A _C filePath	^A _C predicted_labels
1	dbfs:/tmp/flower_photos/sample/100080576_f52e8ee070_n.jpg	> [{"class": "n11939491", "description": "daisy", "probability": 0.8805494}, {"class": "n02219486", "description": "ant", "probability": ...
2	dbfs:/tmp/flower_photos/sample/100930342_92e8746431_n.jpg	> [{"class": "n03930313", "description": "picket_fence", "probability": 0.18473865}, {"class": "n11939491", "description": "daisy", "probability": ...
3	dbfs:/tmp/flower_photos/sample/10140303196_b88d3d6cec.jpg	> [{"class": "n11939491", "description": "daisy", "probability": 0.9535933}, {"class": "n02219486", "description": "ant", "probability": ...

```
df = p_model.transform(image_df)
display(df.select("filePath", (1-p1(df.probability)).alias("p_daisy")))
```

df: pyspark.sql.dataframe.DataFrame

Table

	A ^B _C filePath	1.2 p_daisy
1	dbfs:/tmp/flower_photos/sample/100080576_f52e8ee070_n.jpg	0.968875532556084
2	dbfs:/tmp/flower_photos/sample/100930342_92e8746431_n.jpg	0.13976502414604564
3	dbfs:/tmp/flower_photos/sample/10140303196_b88d3d6cec.jpg	0.9786112803439984

```
from sparkdl import readImages, TFImageTransformer
from sparkdl.transformers import utils
import tensorflow as tf

image_df = readImages(sample_img_dir)

g = tf.Graph()
with g.as_default():
    image_arr = utils.imageInputPlaceholder()
    resized_images = tf.image.resize_images(image_arr, (299, 299))
    # the following step is not necessary for this graph, but can be for graphs with variables, etc
    frozen_graph = utils.stripAndFreezeGraph(g.as_graph_def(add_shapes=True), tf.Session(graph=g), [resized_images])

transformer = TFImageTransformer(inputCol="image", outputCol="transformed_img", graph=frozen_graph,
                                inputTensor=image_arr, outputTensor=resized_images,
                                outputMode="image")
tf_trans_df = transformer.transform(image_df)
```

image_df: pyspark.sql.dataframe.DataFrame = [filePath: string, image: struct]
 tf_trans_df: pyspark.sql.dataframe.DataFrame = [filePath: string, image: struct ... 1 more field]

INFO:tensorflow:Froze 0 variables.
 Converted 0 variables to const ops.
 INFO:tensorflow:Froze 0 variables.
 Converted 0 variables to const ops.



```
from keras.applications import InceptionV3
```

```
model = InceptionV3(weights="imagenet")
```

```
model.save('/tmp/model-full.h5') # saves to the local filesystem
```

```
# move to a permanent place for future use
```

```
dbfs_model_path = 'dbfs:/models/model-full.h5'
```

```
dbutils.fs.cp('file:/tmp/model-full.h5', dbfs_model_path)
```

```
ut[14]: True
```


04 Test



Model Accuracy: The logistic regression model trained using transfer learning achieved an accuracy of 97.1% on the test set.

Prediction Results: Predictions from the pre-trained InceptionV3 model were displayed, showing high probabilities for relevant classes (e.g., "daisy").

Error Analysis: The DataFrame was ordered by the absolute difference between predicted probabilities and actual labels to identify misclassified images.

```
from keras.applications.inception_v3 import preprocess_input
from keras.preprocessing.image import img_to_array, load_img
import numpy as np
from pyspark.sql.types import StringType
from sparkdl import KerasImageFileTransformer

def loadAndPreprocessKerasInceptionV3(uri):
    # this is a typical way to load and prep images in keras
    image = img_to_array(load_img(uri, target_size=(299, 299))) # image dimensions for InceptionV3
    image = np.expand_dims(image, axis=0)
    return preprocess_input(image)

dbutils.fs.cp(dbfs_model_path, 'file:/tmp/model-full-tmp.h5')
transformer = KerasImageFileTransformer(inputCol="uri", outputCol="predictions",
                                       modelFile='/tmp/model-full-tmp.h5', # local file path for model
                                       imageLoader=loadAndPreprocessKerasInceptionV3,
                                       outputMode="vector")

files = ["/dbfs" + str(f.path)[5:] for f in dbutils.fs.ls(sample_img_dir)] # make "local" file paths for images
uri_df = sqlContext.createDataFrame(files, StringType()).toDF("uri")

keras_pred_df = transformer.transform(uri_df)
```

```
└─ uri_df: pyspark.sql.dataframe.DataFrame = [uri: string]
└─ keras_pred_df: pyspark.sql.dataframe.DataFrame
```

```
/databricks/python/local/lib/python2.7/site-packages/keras/models.py:255: UserWarning: No training configuration found in save file: the model was *not* compiled. C[...]  
file it manually.
```

```
warnings.warn('No training configuration found in save file: '
```

```
INFO:tensorflow:Froze 378 variables.
```

```
Converted 378 variables to const ops.
```

```
INFO:tensorflow:Froze 0 variables.
```

```
Converted 0 variables to const ops.
```

```
display(keras_pred_df.select("uri", "predictions"))
```

Table ▾ +



	uri	predictions
1	/dbfs/tmp/flower_photos/sample/100080576_f52e8ee070_n.jpg	> [1,1000,[],[0.00007469155389117077,0.00007630845357198268,0.0001935783657245338,0.000122831101180054...
2	/dbfs/tmp/flower_photos/sample/100930342_92e8746431_n.jpg	> [1,1000,[],[0.0002563267189543694,0.0028356011025607586,0.00012032653467031196,0.0001531501911813393...
3	/dbfs/tmp/flower_photos/sample/10140303196_b88d3d6cec.jpg	> [1,1000,[],[0.00003744077548617497,0.000053084160754224285,0.00009704380499897525,0.0000665588377160...

⬇ 3 rows | 7.68 seconds runtime

Refreshed 2,3

▶ ✓ 1/22/2018 (23s)

```
dbutils.fs.rm(img_dir, recurse=True)
dbutils.fs.rm(dbfs_model_path)
```


Out[17]: True



05

Enhancements

Can we get better result?



Enhancement

Future enhancements planned for Deep Learning Pipelines include:

1. **Distributed Hyper-parameter Tuning:** Leveraging Spark MLlib Pipelines to perform hyper-parameter tuning in a distributed manner.
2. **SQL Functions:** Deploying deep learning models as SQL functions to make them easily accessible via SQL queries.
3. **Enhanced Model Support:** Adding support for more models and improving the integration with other deep learning libraries.

06

Conclusion



- 
- Deep Learning Pipelines offers a comprehensive suite of tools to integrate deep learning with Apache Spark, making it accessible for scalable data processing and model deployment.
 - Its high-level APIs simplify the application of deep learning models and transfer learning, while future enhancements promise to further extend its capabilities, making it a valuable resource for data scientists and engineers working with big data and deep learning.
- 
- 




07

References



Introducing Deep Learning Pipelines for Apache Spark

GitHub



Thanks!

CREDITS: This presentation template was created by [Slidesgo](#), and includes icons by [Flaticon](#), and infographics & images by [Freepik](#)

Please keep this slide for attribution

